



Enabling interactive data exploration

Yandex

What is ClickHouse?

- RDBMS for analytics
 - SQL
 - FOSS
- Distributed
 - Cross-datacenter replication
 - Tolerant to a single-datacenter failure
- Linear scaling
 - 100s of servers
 - 100G-1000G records daily
- Blazing fast
 - interactive data exploration

These slides have a lot of [links!](#)

Download them at <insert link and code>



History

- Yandex.Metrica (think Google Analytics)
 - 30B rows daily, 3PB total, 700 machines
 - Processing speed up to 2TB/s
- Used a MyISAM solution with pre-aggregation
 - Didn't scale for the growing load
 - Could only build pre-defined reports
- Started developing own columnar DB in 2009
- Enables “double real time” reporting
 - Add new events in real time
 - Build custom reports in real time
- Becomes the main engine in 2012
- Open-sourced in 2016

Adoption

- Yandex-wide
 - business analytics
- Thousands of companies worldwide
 - [analyzing 1M DNS queries per second](#) (Cloudflare)
 - [geospatial processing](#) (Carto)
 - [real-time ad analytics platform](#) (LifeStreet)
 - storage performance analysis and monitoring (Infinidat)
 - AdTech, FinTech, sensors, logs, event streams, time series, etc.

Unusual applications:

- Blockchain transaction history
 - blockchair.com
 - bloxy.info
- CERN LHCb experiment
- Bioinformatics

When to use

- stream of well-structured, immutable events
 - (mostly) fixed schema
 - append-only
 - heavy-weight ALTER UPDATE/DELETE as a “GDPR escape hatch”
- flexible real-time reporting
 - queries finish in seconds, not hours
 - no preprocessing needed
 - enables ad-hoc experiments

When NOT to use

- OLTP
 - no transactions
 - single INSERT is atomic, but no cross-query atomicity
 - very heavy UPDATES
- key/value storage
 - sparse indexes
 - not suitable for point reads
- document/blob storage
 - optimized for records < 100 kB
- highly normalized data
 - optimized for star schema

How fast?

[1.1 Billion Taxi rides benchmark](#): 1.1G records, 51 columns, 500 GB uncompressed CSV

Query 1	Query 2	Query 3	Query 4	Setup
0.005	0.01	0.10	0.188	BrytlytDB 2.1 & 5-node IBM Minsky cluster
0.051	0.15	0.05	0.794	kdb+/q & 4 Intel Xeon Phi 7210 CPUs
0.241	0.83	1.21	1.781	ClickHouse, 3 x c5d.9xlarge cluster
0.762	2.47	4.13	6.041	BrytlytDB 1.0 & 2-node p2.16xlarge cluster
1.034	3.06	5.35	12.748	ClickHouse, Intel Core i5 4670K
1.56	1.25	2.25	2.97	Redshift, 6-node ds2.8xlarge cluster
2	2.00	1.00	3	BigQuery
2.362	3.56	4.02	20.412	Spark 2.4 & 21 x m3.xlarge HDFS cluster
14.389	32.15	33.45	67.312	Vertica, Intel Core i5 4670K

Why so fast?

Read less data

- Data locality
 - columnar storage
 - sorted by PK
- Compression

Process data faster

- Parallelism
 - multithreading
 - distributed queries
- Efficient computation
 - 17 different GROUP BY algorithms
 - vectorized query execution with SIMD

Deployment

- Repos for major Linux distros
- Docker containers
- One self-contained binary
 - works everywhere
- ZooKeeper if you need replication
- Test it on your laptop
 - runs with minimal resources
- Stable release every two weeks
 - Thousands of scenarios tested for each change
- No data migration on update
 - Just run a new server

Data ingestion

- INSERTs
- Many [data formats](#)
 - CSV, JSON, Parquet, CapnProto, ORC, ...
- Batching for optimal performance
 - [Buffer](#) table engine
 - [Kafka](#) table engine
 - [Third-party solutions](#) — chproxy, kittenhouse etc.

Analyzing data

- A rich SQL dialect
 - [strong typing](#)
 - [higher order functions](#) like *arrayMap*
 - variety of [aggregate function](#) —
quantiles, cardinality estimators etc.
 - [sampling](#)
 - [Nested](#) type for key/value records
 - [LowCardinality](#) type for dictionary
encoding
- BI tools support
 - [Holistics](#) (Singapore-based)
 - [Tableau](#)
 - [Apache Superset](#)
 - others via ODBC/SQLAlchemy/...

clickhouse-local

The full power of ClickHouse engine over a data file.

```
$ clickhouse local
  --file ~/hits_v1.tsv
  --structure 'WatchID UInt64,  JavaEnable UInt8, ...'
  --query 'SELECT UserID, SearchPhrase, count() FROM table GROUP BY UserID,
SearchPhrase'
```

Read 8873898 rows, 7.88 GiB in 5.208 sec., 1704038 rows/sec., **1.51 GiB/sec.**

UserID	SearchPhrase	count()
8410854169855355129	пальные кость играть терхи	3

Interfacing with other systems

Connect to ClickHouse

- Native binary protocol
 - Drivers for Python, Go, C++, ...
- RESTful [HTTP](#)
- [ODBC](#)
- [JDBC](#)
- MySQL wire protocol
- PostgreSQL
 - [clickhouse_fdw](#)
 - [pg2ch](#) (logical replication)

Connect from ClickHouse

- [File](#)
- [HDFS](#)
- [URL](#)
- [MySQL](#)
- [ODBC](#)
- [External dictionaries](#)

Getting help

- Check the [docs at our site](#)
- Create issues on [github](#)
- Ask on [Stack Overflow](#)
- Email us at clickhouse-feedback@yandex-team.ru
- Join [English](#) and [Russian](#) chats in Telegram
- Get commercial support from [Altinity](#) and others
- [... and more](#)

Thank you!

Questions?

- <https://clickhouse.yandex>
- <https://github.com/ClickHouse/ClickHouse>
- clickhouse-feedback@yandex-team.ru

