



# Altinity



## ClickHouse on Kubernetes!

Alexander Zaitsev, Altinity  
Limassol, May 7th 2019



## Altinity Background

- Premier provider of software and services for ClickHouse
- Incorporated in UK with distributed team in US/Canada/Europe
- US/Europe sponsor of ClickHouse community
- Offerings:
  - 24x7 support for ClickHouse deployments
  - Software (Kubernetes, cluster manager, tools & utilities)
  - POCs/Training

# What is Kubernetes?

**“Kubernetes is the new Linux”**

Actually it's an open-source platform to:

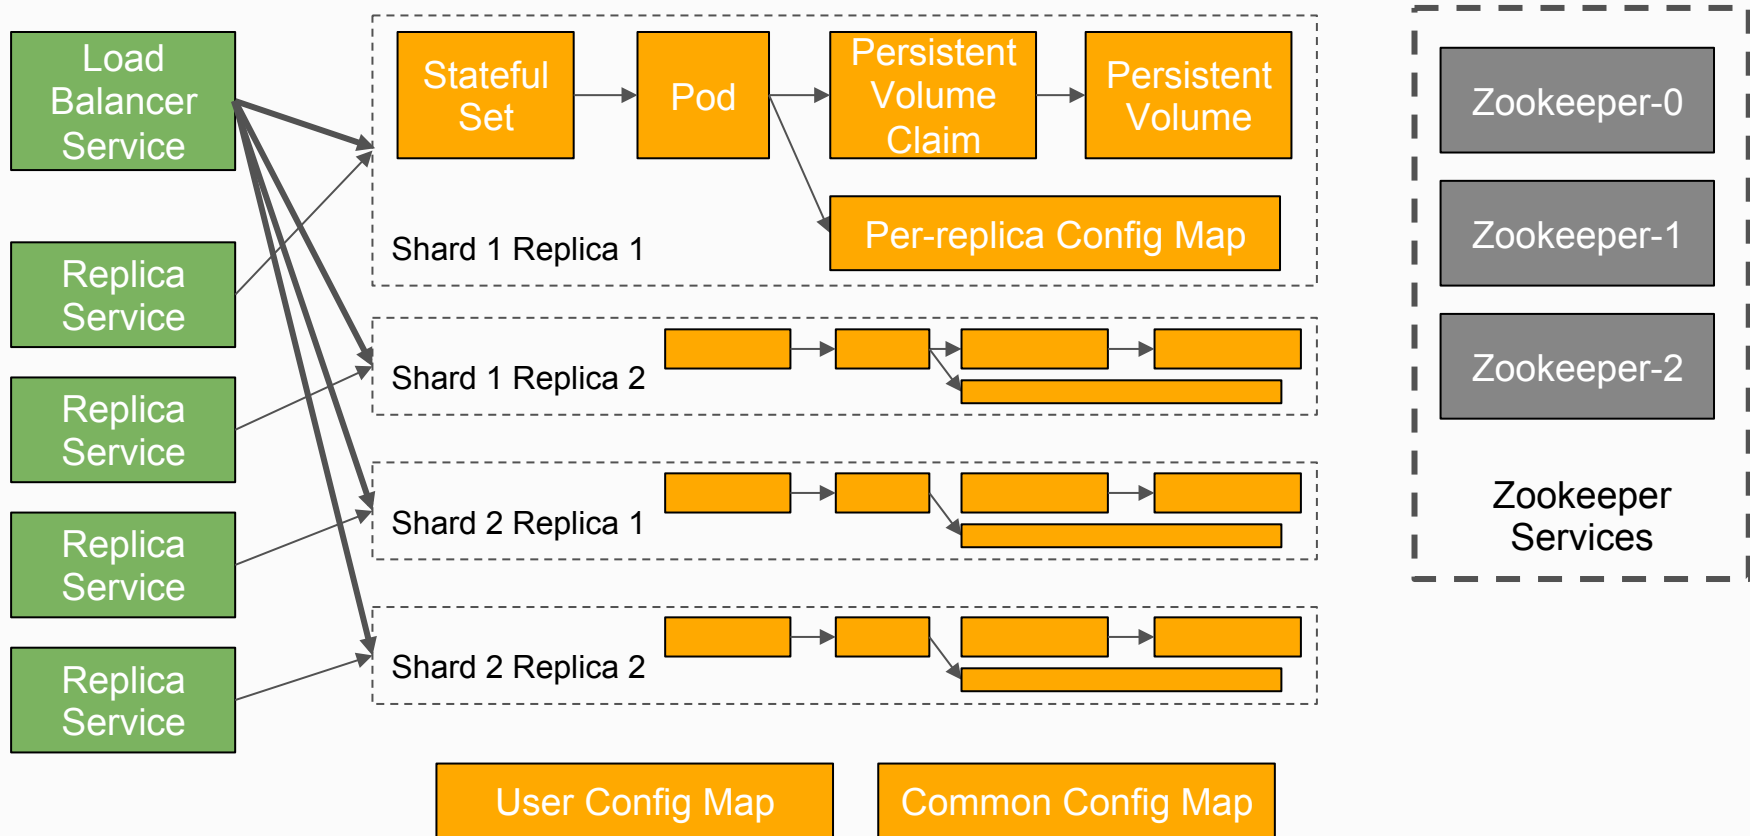
- manage container-based systems
- build distributed applications declaratively
- allocate machine resources efficiently
- automate application deployment



# Why run ClickHouse on Kubernetes?

1. Other applications are already there
2. Portability
3. Bring up data warehouses quickly
4. Easier to manage than deployment on hosts

# What does ClickHouse look like on Kubernetes?

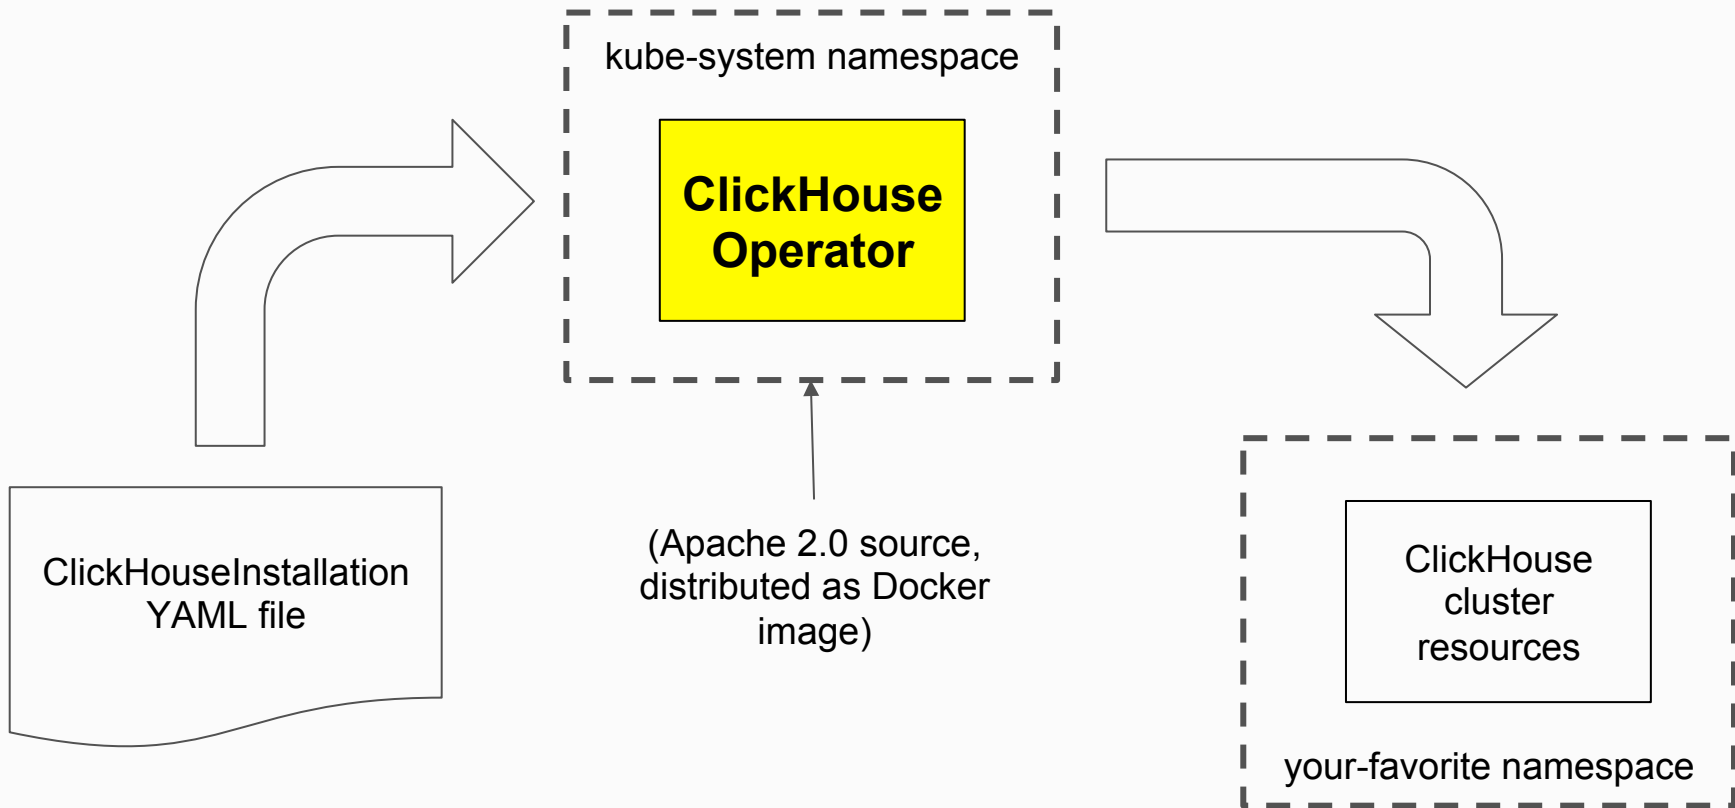


# Challenges running ClickHouse on Kubernetes?

1. Provisioning
2. Persistence
3. Networking
4. Transparency



# The ClickHouse operator turns complex data warehouse configuration into a single easy-to-manage resource



# Altinity ClickHouse Operator Quick Start



# Installing and removing the ClickHouse operator

[Optional] Get sample files from github repo:

```
git clone https://github.com/Altinity/clickhouse-operator
```

Install the operator:

```
kubectl apply -f clickhouse-operator-install.yaml
```

or:

```
kubectl apply -f https://raw.githubusercontent.com/Altinity/clickhouse-operator/master/manifests/operator/clickhouse-operator-install.yaml
```

# Let's start with a single-node cluster

```
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "demo-01"
spec:
  configuration:
    clusters:
      - name: "demo-01"
        layout:
          type: Standard
          shardsCount: 1
          replicasCount: 1
```

**WARNING: This installation  
lacks persistent storage**

**See examples in later slides for  
storage definition**

# kubectl is our tool

```
$ kubectl create namespace demo
namespace/demo created
```

```
$ kubectl apply -f docs/examples/demo-01.yaml -n demo
clickhouseinstallation.clickhouse.altinity.com/demo-01 created
```

```
$ kubectl get all -n demo
```

NAME	READY	STATUS	RESTARTS	AGE
pod/chi-a82946-2946-0-0-0	1/1	Running	0	52s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/chi-a82946-2946-0-0	ClusterIP	None	<none>	8123/TCP,9000/TCP,9009/TCP
52s				
service/clickhouse-demo-01	LoadBalancer	10.108.198.248	<pending>	8123:31609/TCP,9000:32086/TCP
52s				

NAME	DESIRED	CURRENT	AGE
statefulset.apps/chi-a82946-2946-0-0	1	1	52s

## Next let's add a shard

```
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "demo-01"
spec:
  configuration:
    clusters:
      - name: "demo-01"
        layout:
          type: Standard
          shardsCount: 2
          replicasCount: 1
```

```
$ kubectl apply -f docs/examples/demo-01.yaml -n demo
clickhouseinstallation.clickhouse.altinity.com/demo-01 configured
```

# How to access your ClickHouse data warehouse on Kubernetes

## Connect from within Kubernetes using service DNS name

```
# Use load balancer
clickhouse-client --host clickhouse-demo-01.test
# Connect to specific node
clickhouse-client --host chi-a82946-2946-0-0.test
# Connect via pod entry
kubectl exec -it chi-a82946-2946-0-0-0 -n demo -- clickhouse-client
```

## Connect from outside Kubernetes using Ingress or Nodeport

```
# Kops deployment on AWS configures external ingress.
clickhouse-client --host $AWS_ELB_HOST_NAME
```

## Replication requires Zookeeper to be enabled

Install minimal Zookeeper in separate namespace.

```
kubectl create ns zoons  
kubectl apply -f zookeeper-1-node.yaml -n zoons  
watch kubectl -n zoons get all
```

Note ZK node DNS name: **zookeeper-0.zookeepers.zoons**

You can also install using helm \*or\* use external ZK cluster

## After inserting a 'zookeepers' clause we can add replicas

```
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "demo-01"
spec:
  configuration:
    zookeeper:
      nodes:
        - host: zookeeper-0.zookeepers.zoons
          port: 2181
  clusters:
    - name: "demo-01"
      layout:
        type: Standard
        shardsCount: 2
        replicasCount: 2
```

**TIP: Confirm the DNS name of Zookeeper from within a pod**

**NOTE: Non-replicated tables do not replicate automatically when replicas are added**

# We can add and modify users with the 'users' clause

```
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "demo-01"
spec:
  configuration:
    users:
      demo/default: secret
      demo/password: demo
      demo/profile: default
      demo/quota: default
      demo/networks/ip: "::/0"
  clusters:
    - name: "demo-01"
      layout:
        type: Standard
        shardsCount: 2
        replicasCount: 1
```

**TIP: User and profile changes take a few minutes to propagate. Confirm changes using clickhouse-client**



# To make storage persistent and set properties add an explicit volume claim template with class and size

```
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "storage"
spec:
  defaults:
    deployment:
      volumeClaimTemplate: storage-vc-template
  templates:
    volumeClaimTemplates:
      - name: storage-vc-template
        persistentVolumeClaim:
          spec:
            accessModes:
              - ReadWriteOnce
            resources:
              requests:
                storage: 2Gi
  configuration:
```

**TIP: Check syntax carefully as errors may result in failures to allocate or mount volumes**

**TIP: Confirm storage by 'kubectl exec' into pod; run 'df -h' to confirm mount**

storageClassName can be used to set the proper class of storage as well as disable dynamic provisioning

Use kubectl to find available storage classes:

```
kubectl describe StorageClass
```

Bind to default storage:

```
spec:  
  storageClassName: default
```

Bind to gp2 type

```
spec:  
  storageClassName: gp2
```

Disable dynamic provisioning and use static PVs:

```
spec:  
  storageClassName: "
```

# Set the ClickHouse version using a podTemplate

```
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "demo-02"
spec:
  defaults:
    deployment:
      podTemplate: clickhouse-stable
      volumeClaimTemplate: storage-vc-template
  templates:
    podTemplates:
      - name: clickhouse-stable
        containers:
          - name: clickhouse
            image: yandex/clickhouse-server:19.4.3.11
        volumeClaimTemplates:
# Etc.
```

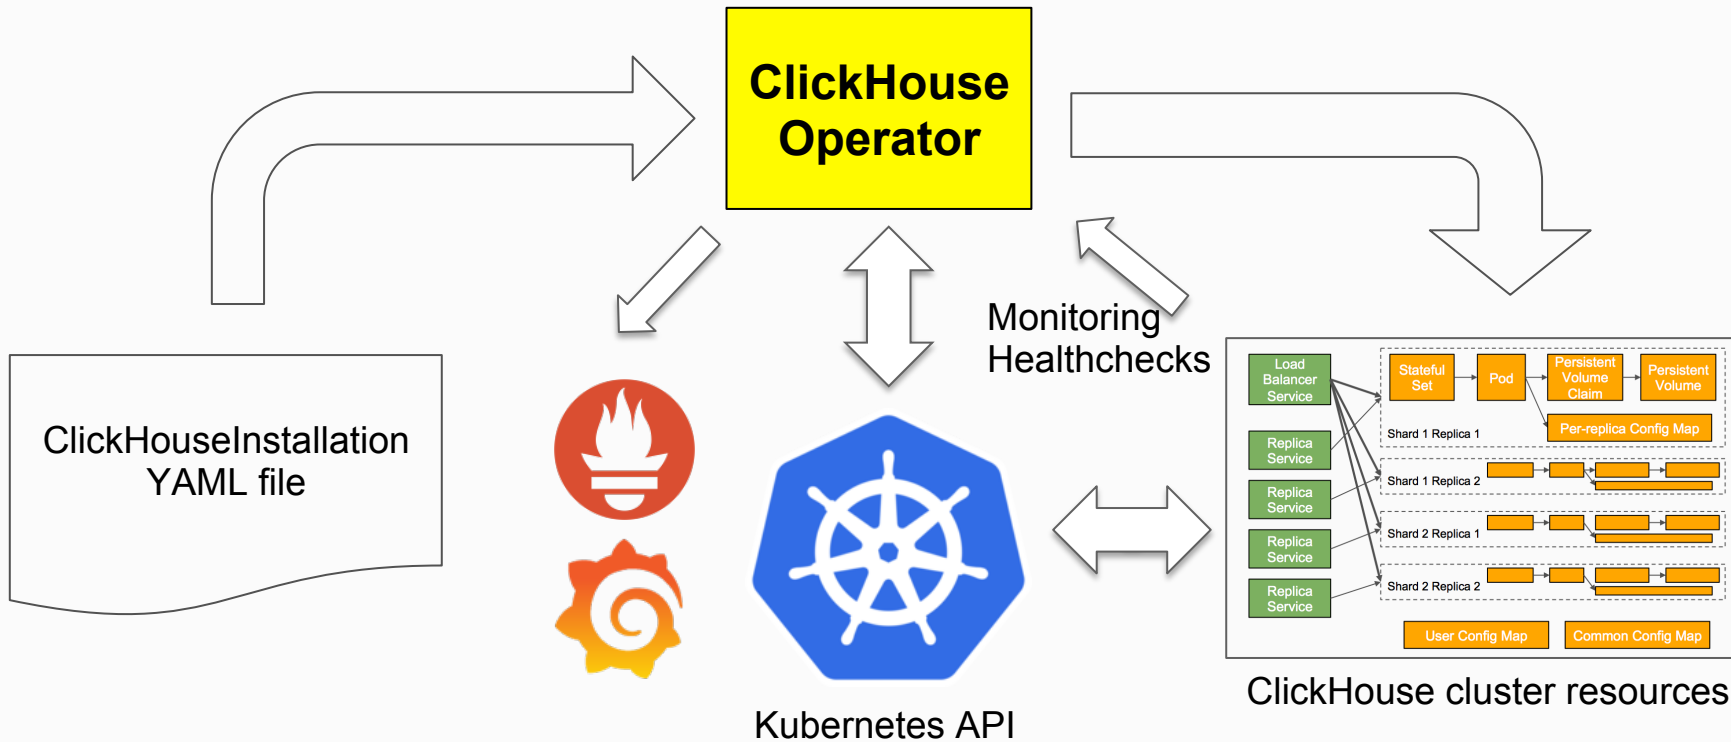
**TIP: Always specify the image version fully; do not use 'latest' tag**

# More pod template tricks: controlling resources

```
spec:
  defaults:
    deployment:
      podTemplate: clickhouse-stable
      volumeClaimTemplate: storage-vc-template
  templates:
    podTemplates:
      - name: clickhouse-stable
        containers:
          - name: clickhouse
            image: yandex/clickhouse-server:19.4.3.11
            resources:
              requests:
                memory: "512Mi"
                cpu: "500m"
              limits:
                memory: "512Mi"
                cpu: "500m"
```

# Etc.

# Operator = deployment + monitoring + operation




# Codified Operational Knowledge


- Available already:
  - Monitoring to Prometheus / Grafana
  - Automatically create schema when adding shards/replicas
  - Track IP address changes when pod is re-provisioned
  - Rollback to the previous configuration in case of failure
- Planned / in-progress:
  - Configurable defaults
  - Multi-zone deployments
  - Advanced health checks
  - Automatic / semi-manual recovery
  - Replica re-provisioning
  - Storage management (tiered storage, backups)



# Performance. AWS vs k8s@AWS. 1 node



# Performance. AWS vs k8s@AWS. 3 nodes





# Kubernetes Container Network Interface (CNI)

- Calico
- Canal (Flannel + Calico)
- flannel
- kopeio-vxlan
- kube-router
- romana
- Weave Net

To be tested!

## Other Roadmap features

- Store history of configuration changes in a separate resource
- 'Canary testing' of configuration changes at clean ClickHouse instance
- Operator status transparency
- Integration with ZooKeeper operator for automatic ZK provisioning
- Default configuration templates, including networking
- Integration with Altinity Cluster Manager

## Advice, encouragement, and caveats

- Clickhouse operator is in beta
- Operator does not always detect errors in manifest files
- Error logging is limited, will be improved shortly
- Connectivity is a work in progress
- It's a great way to explore cluster configurations
- Kubernetes is fun!

**Please explore the operator and log issues on Github!!!**

## More information on Altinity ClickHouse Operator...

ClickHouse Operator Github Project:

<https://github.com/Altinity/clickhouse-operator>

Altinity Blog -- <https://www.altinity.com/blog>

Webinars -- <https://www.altinity.com/webinarspage>

# Questions?

# Thank you!

Contacts:

[info@altinity.com](mailto:info@altinity.com)

Visit us at:

<https://www.altinity.com>

Read Our Blog:

<https://www.altinity.com/blog>