




A large, abstract graphic of a blue wave with white speckles, resembling a data visualization or a stylized ocean, spans the middle of the slide. The wave is composed of many thin, overlapping lines and dots, creating a sense of movement and complexity.



# Unusual applications of a column-oriented database

Sascha Dienelt – Senior Software Developer (data warehouse, bid management)

# Usual applications of column-oriented DBMS


- Fetch a huge amount of time series data
- Aggregate many rows at once
- Process date ranges






# What we need to do


- Normalize values




- Filter duplicates




- Build chains



# What we need to do



# Different ClickHouse clusters



- cluster 1: raw data
- cluster 2: processed information
- Solution: remote table functions (ClickHouse <-> ClickHouse)
- Communication overhead → don't use too extensive
- Alternative: Distributed table engine


```
INSERT INTO processed.events
    (eventId, insertTimestamp, typeId)
SELECT
    eventId,
    insertTimestamp,
    dictGetUInt16('Global_MasterData.Event_Types', 'typeId', visitParamExtractString(eventData, 'typeName')) AS typeId
FROM remoteSecure(
    'raw-clickhouse-hostname.diva-e.com',
    'rawSchemaName',
    'EventTableName',
    'clickHouseUserName',
    'sup3rS3<ur3Passw0rd'
) AS rawEvents
WHERE insertTimestamp BETWEEN '2019-09-01 00:00:00' AND '2019-09-02 00:00:00';
```

# Master data and further processing in MySQL



- Master data must be kept in a RDBMS (here: MySQL)
- Implementation of very complex logic is using MySQL
- Solution: MySQL table engine (ClickHouse <-> MySQL)

```
CREATE TABLE default.tmp_MySqlEventConnector_4815162342 (  
    eventId UInt64,                -- UNSIGNED BIGINT NOT NULL  
    insertTimestamp DateTime,      -- DATETIME NOT NULL  
    typeId Nullable(UInt8)        -- UNSIGNED TINYINT DEFAULT NULL  
) ENGINE = MySQL(  
    'mysql-hostname.diva-e.com',  
    'schemaName',  
    'tableName',  
    'mySqlHouseUserName',  
    'an0th3rSup3rS3<ur3Passw0rd'  
)  
;  
  
INSERT INTO default.tmp_MySqlEventConnector_4815162342 (eventId, insertTimestamp, typeId)  
SELECT eventId, insertTimestamp, typeId FROM processed.events  
WHERE insertTimestamp BETWEEN '2019-09-01 00:00:00' AND '2019-09-02 00:00:00';  
  
DROP TABLE default.tmp_MySqlEventConnector_4815162342;
```



# Client settings require different queries

- processing many elements at once → fast
- many “small” queries → very slow
- Clients want different filters and chaining rules
- Solution: combine many steps into one query

```
SELECT eventId FROM processed.events
WHERE clientId = 1
      AND insertTimestamp >= now() - INTERVAL 14 DAY
      AND typeId IN (1, 2, 3);
```

```
SELECT eventId FROM processed.events
WHERE clientId = 2
      AND insertTimestamp >= now() - INTERVAL 30 DAY
      AND typeId IN (2, 3);
```

```
SELECT eventId FROM processed.events
WHERE clientId = 3
      AND insertTimestamp >= now() - INTERVAL 7 DAY
      AND typeId IN (1);
```

```
SELECT eventId FROM processed.events
WHERE insertTimestamp >= now() - INTERVAL multiIf(
      clientId = 1, 14,
      clientId = 2, 30,
      clientId = 3, 7,
      10
) DAY
AND multiIf(
      clientId = 1 AND typeId IN (1, 2, 3), 1,
      clientId = 2 AND typeId IN (2, 3), 1,
      clientId = 3 AND typeId IN (1), 1,
      0
);
```







# High memory usage and long run times

- Lot of rows need to be covered
- Many complex joins
- Solutions:
  - Use pseudo temporary tables
    - Pro:
      - simpler queries
      - save memory
      - reusable calculations
    - Con:
      - more overhead
      - cleanup required



```
CREATE TABLE default.tmp_IntermediateResult_4815162342 (  
  eventId FixedString(15),  
  chainStartEventId FixedString(15),  
  chainEndEventId FixedString(15)  
) ENGINE = Log;
```

```
INSERT INTO default.tmp_IntermediateResult_4815162342  
(eventId, chainStartEventId, chainEndEventId)  
SELECT conversions.eventId,  
       argMin(eventId, insertTimestamp) AS chainStartEventId,  
       argMax(eventId, insertTimestamp) AS chainEndEventId  
FROM processed.events AS conversions  
INNER JOIN processed.events AS clicks ON (clicks.typeId = 2  
      AND clicks.sessionId = conversions.sessionId)  
WHERE conversions.typeId = 1;  
  
DROP TABLE default.tmp_IntermediateResult_4815162342;
```



# High memory usage and long run times

- Solutions:

- Use clever partitioning 
- Filter by sorting key 

```
CREATE TABLE raw.events (  
  eventId FixedString(15),  
  insertTimestamp DateTime,  
  sessionId FixedString(18),  
  clientId UInt32,  
  eventType Enum8('click' = 1, 'conversion' = 2),  
  eventData String  
) ENGINE = MergeTree()  
PARTITION BY toYYYYMM(insertTimestamp)  
ORDER BY (sessionId, typeId, insertTimestamp, eventId);
```

```
SELECT sessionId,  
  argMin(eventId, insertTimestamp) AS firstEventId,  
  argMax(eventId, insertTimestamp) AS lastEventId  
WHERE insertTimestamp >= '2019-08-01 00:00:00'  
  AND insertTimestamp < '2019-09-01 00:00:00'  
  AND typeId = 2  
GROUP BY sessionId;
```




# High memory usage and long run times

- Solutions:
  - Do not join plain table, join pre-filtered subselect

```
SELECT key, h.column1, l.column2
FROM default.hugeTable AS h
INNER JOIN default.veryLargeTable AS l USING (key)
WHERE h.insertDate BETWEEN '2019-01-01' AND '2019-01-31'
      AND l.insertDate BETWEEN '2019-02-01' AND '2019-02-28';
```

```
SELECT key, h.column1, l.column2
FROM (
  SELECT key, column1
  FROM default.hugeTable
  WHERE insertDate BETWEEN '2019-01-01' AND '2019-01-31'
) AS s
INNER JOIN (
  SELECT key, column2
  FROM default.veryLargeTable
  WHERE insertDate BETWEEN '2019-02-01' AND '2019-02-28'
) AS l USING (key);
```



## Before and after

- Complex process
  - ClickHouse exports data
  - TSV files in NFS folders
  - MySQL imports data
  - Processed data is copied to ClickHouse
- Simpler process
  - No file usage anymore
  - Most processing directly in ClickHouse
  - Less copy overhead
  - MySQL: less load and storage
- Easier to maintain
- Faster low-level reporting

## Questions?

Copyright © diva-e

Alle Angaben basieren auf dem derzeitigen Kenntnisstand. Änderungen vorbehalten. Dieses Dokument der diva-e Digital Value Excellence GmbH ist ausschließlich für den Adressaten bzw. Auftraggeber bestimmt. Es bleibt bis zur einer ausdrücklichen Übertragung von Nutzungsrechten Eigentum der diva-e. Jede Bearbeitung, Verwertung, Vervielfältigung und/oder gewerbsmäßige Verbreitung des Werkes ist nur mit Einverständnis von diva-e zulässig.

*All content is based on the current state of communication. Subject to change. This document of diva-e Digital Value Excellence GmbH is only intended for the client. It belongs to diva-e until its explicit transfer of usage rights. Any adaptation, utilization, copy and/or professional spreading has to be approved by diva-e.*