

# 10 Good Reasons to Use



ClickHouse

*by Ramazan Polat*

## Who I am



*Ramazan Polat*

Former **DBA**, currently **Software Architect** at 

**SGK**(Social Security Institution) is the largest\*  
government institution of Turkey in terms of *in-house*  
*generated data size and transactions per second.*

*\* Possibly, not sure about it*

# Why I am here


I've been using a number of database management systems for 15+ years, but ended up always using ClickHouse, even for cases where ClickHouse doesn't perfectly fit.

## *Why?*

I have a lot of reasons, here is **10** of them, so you can use ClickHouse too.

# Remember these

1) Data is growing faster than computing power\*



2) Unix timestamp is currently over 1.5 Billion (in fact, as of writing this, it is 1574114808)

\* <https://www.newsweek.com/2014/08/15/computers-need-be-more-human-brains-262504.html>

# #1 Speed

- ✓ *Inserts are instant!*
- ✓ *Selects are blazing fast!*
- ✓ *Handling billions of rows sub-second!*



# #2 Scalability




*Uses to all CPU cores in single machine*



Thanks to **vectorized execution** and **parallel processing**

## #2 Scalability

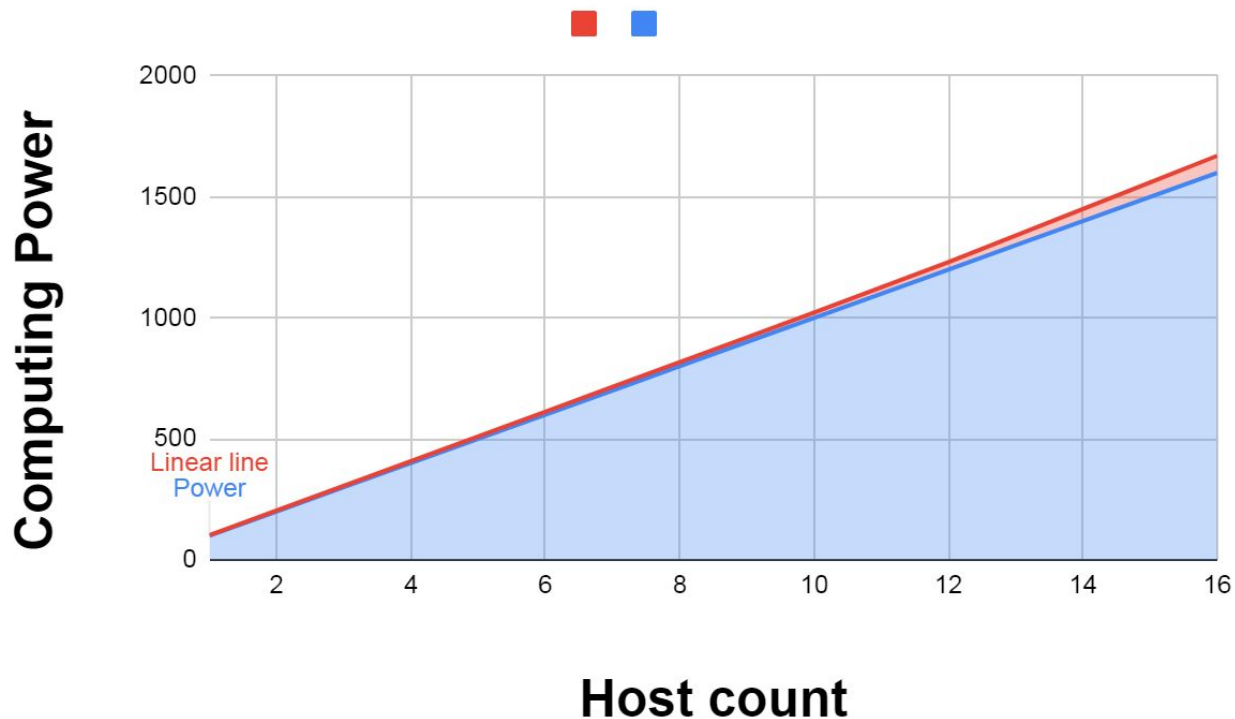
✓ *Scales horizontally across multiple hosts*



*From clients, it looks like one giant database*

## #2 Scalability

✓ *Linearly scaling.*





# #3 Compression

- ✓ *Compression is column based*
- ✓ *Encoding helps data compress better*
- ✓ *Ritch compression options*
- ✓ *Custom compression for different type of data*

# #3 Compression

*Encoding maps data in a different bit layout*

Encodings	
<b>LowCardinality</b>	String with a few values(up to 10K)
<b>Delta</b>	Difference between consecutive values
<b>DoubleDelta</b>	Difference between consecutive deltas, good for slowly changing sequences
<b>Gorilla</b>	Efficient for values that does not change often
<b>T64</b>	Strips lower and higher bits that does not change, good for big numbers in a small range

# #3 Compression


*There is a trade-off between speed and size*

Compression algos	
<b>LZ4</b>	Faster compression with smaller compression ratio
<b>ZSTD</b>	Slower compression but compresses better

```
CREATE TABLE test (  
  city LowCardinality(String),  
  speed UInt32 CODEC(Delta),  
  acceleration UInt32 CODEC(DoubleDelta),  
  humidity UInt32 CODEC(Gorilla, LZ4),  
  year UInt32 CODEC(T64, ZSTD)  
) ENGINE = ...
```

# #4 Production Ready

✓ *Fault tolerant*

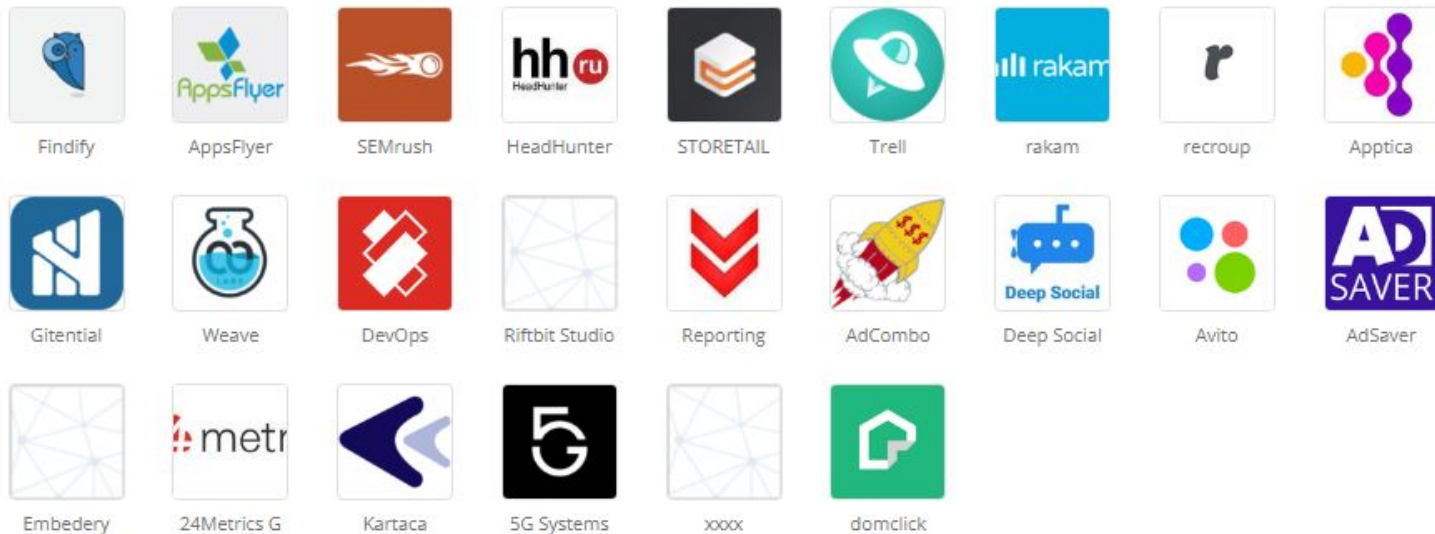


*Multi-master replication*

# #4 Production Ready



*Used by several companies worldwide*



*This list is taken from [stackshare.io/clickhouse](https://stackshare.io/clickhouse)*

# #4 Production Ready

✓ *Already being used in Yandex Metrica\**

A cluster of **374** servers  
More than **13 trillion** rows  
**20 billion** events per day  
**17 PB** data(**2 PB** compressed)

**Just a reminder:** *The unixtimestamp is number of seconds that have elapsed since the 1 January 1970, which is around **1.5 Billion***

*\* Yandex Metrica is World's 2nd largest analytics database after Google Analytics*

# #5 Integration

- ✓ Connects to any other JDBC database and uses their tables as ClickHouse table



```
SELECT *  
FROM jdbc('mysql://localhost:3306/?user=root&password=root', 'schema', 'table')
```

This means **any JDBC database table** is also a **ClickHouse table**

## #5 Integration



*Even can connect to ClickHouse itself without any configuration*

```
SELECT * FROM remote('ch2', 'db_name', 'table_name', 'user', 'password')
```



*Connects to REST services!*

```
SELECT * FROM url('http://127.0.0.1:12345/', CSV, 'column1 String, column2 UInt32')
```



# #5 Integration

✓ *Clickhouse integrates into anything\**



*... and much more*

\* almost

# #6 Awesome Table Capabilities

- ✓ *Enums*
- ✓ *Partitioning*
- ✓ *Arrays*
- ✓ *Nested columns*
- ✓ *TTL tables*
- ✓ *Materialized views*
- ✓ *Live views*

# #6 Awesome Table Capabilities

## *Enums*

```
CREATE TABLE enum_table(  
    `day` Enum8('SUNDAY' = 0, 'MONDAY' = 1, 'TUESDAY' = 2,  
        'WEDNESDAY' = 3, 'THURSDAY' = 4, 'FRIDAY' = 5, 'SATURDAY' = 6)  
)  
ENGINE = ...
```

```
INSERT INTO enum_table VALUES('SATURDAY'),('SUNDAY')
```

```
SELECT *  
FROM enum_table
```

day
SATURDAY
SUNDAY

# #6 Awesome Table Capabilities

## *Enums*

```
CREATE TABLE enum_table(  
    `day` Enum8('SUNDAY' = 0, 'MONDAY' = 1, 'TUESDAY' = 2,  
        'WEDNESDAY' = 3, 'THURSDAY' = 4, 'FRIDAY' = 5, 'SATURDAY' = 6)  
)  
ENGINE = ...
```

```
INSERT INTO enum_table VALUES('SATURDAY'),('SUNDAY')
```

```
SELECT *  
FROM enum_table
```

<pre>day SATURDAY SUNDAY</pre>
--

```
SELECT CASE day  
    WHEN 0 THEN 'SUNDAY'  
    WHEN 1 THEN 'MONDAY'  
    WHEN 2 THEN 'TUESDAY'  
    WHEN 3 THEN 'WEDNESDAY'  
    WHEN 4 THEN 'THURSDAY'  
    WHEN 5 THEN 'FRIDAY'  
    WHEN 6 THEN 'SATURDAY'  
END FROM enum_table;
```

# #6 Awesome Table Capabilities

## Partitioning

- Each partition is stored separately in order to simplify manipulations of this data
- Each partition can be detached, attached or dropped instantly

```
CREATE TABLE visits(  
  VisitDate Date,  
  Hour UInt8,  
  ClientID UUID  
) ENGINE = MergeTree()  
PARTITION BY toYYYYMM(VisitDate)  
ORDER BY Hour;
```

```
PARTITION BY (toMonday(StartDate), EventType)
```

```
ALTER TABLE visits DETACH PARTITION 201910;  
ALTER TABLE visits ATTACH PARTITION 201911;  
ALTER TABLE visits DROP PARTITION 201911;
```

Partitions	VisitDate	Hour	ClientID
Part_1	2019-10-11	14	e0d6e0ff
	2019-10-12	16	d2af7d50
	2019-10-28	17	1aef1ff5
Part_2	2019-11-01	18	0f4def0b
	2019-11-17	19	46638b95
	2019-11-21	21	d6e0af7d
	2019-11-23	23	38b9ef1f

# #6 Awesome Table Capabilities

## Arrays

- Columns can be array of any type
- Great flexibility

```
CREATE TABLE timeseries(  
  entity String,  
  ts UInt64,  
  m Array(String),  
  v Array(Float32),  
  d Date DEFAULT toDate(ts)  
) ENGINE = MergeTree(d, ts, 8192)
```

entity	ts	m	v
cpu1	1574018595	[temp,load]	[78, 0.85]
cpu2	1574018674	[load]	[0.44]
cpu7	1574019333	[ghz, temp]	[3.1, 0.4]
cpu4	1574019501	[load, ghz]	[0.9, 3.8]

# #6 Awesome Table Capabilities

## *Nested Data*

```
CREATE TABLE customers(  
  custId UInt32,  
  name String,  
  Orders Nested(  
    date Date,  
    items UInt16,  
    price Float32)  
) ENGINE = TinyLog()
```

custId	name	Orders		
		date	items	price
123	Joe Lee	2019-01-23	3	€ 34.56
234	Kate Hall	2019-01-24	2	€ 23.45
		2019-01-25	4	€ 21.09
456	Ann Cook	2019-01-26	5	€ 18.91

```
INSERT INTO customers VALUES (123, 'Joe Lee', ['2019-01-23'], [3], [34.56])
```

# #6 Awesome Table Capabilities

## ***TTL Tables***

*Automatically deletes rows based on a conditions happens on row data*

```
CREATE TABLE traffic2(  
    datetime DateTime,  
    custId UInt32  
) ENGINE = MergeTree  
PARTITION BY toYYYYMM(datetime)  
ORDER BY (datetime, custId)  
TTL datetime + INTERVAL 6 MONTH;
```

*Deletes data after 6 months*



# #6 Awesome Table Capabilities

## Materialized Views

*Automatically aggregates data on inserts*

```
CREATE TABLE CCEvents2(  
  dt DateTime,  
  custId UInt32,  
  name String,  
  ccNumber String,  
  posId UInt32,  
  purchase Float32  
) ENGINE = MergeTree  
ORDER BY (dt, custId);
```

```
SELECT *  
FROM CCEvents
```

dt	custId	name	ccNumber	posId	purchase
2019-10-11 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	557	98.9
2019-10-17 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	345	145.3
2019-10-18 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	124	17.4
2019-10-19 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	548	51
2019-11-01 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	61.1
2019-11-01 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	61.1
2019-11-13 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	852	45.3
2019-11-14 21:06:49	123	Joe	XXXX-XXXX-XXXX-XXXX	499	245.6
2019-11-15 21:06:49	123	Joe	XXXX-XXXX-XXXX-XXXX	756	5.8
2019-11-21 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	83.4
2019-12-07 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	20.2
2019-12-27 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	21.3

# #6 Awesome Table Capabilities

## Materialized Views

*Automatically aggregates data on inserts*

```
SELECT
    name,
    toYYYYMM(dt) AS month,
    sum(purchase)
FROM CCEvents
GROUP BY
    name,
    month
```

name	month	sum(purchase)
Joe	201911	296.70000553131104
Kate	201910	98.9000015258789
Kate	201911	205.5999984741211
Joe	201910	213.7000026702881
Kate	201912	41.5

5 rows in set. Elapsed: 0.010 sec.

```
CREATE MATERIALIZED VIEW CCMonthlySums
ENGINE = SummingMergeTree
ORDER BY (name, month)
POPULATE AS
SELECT
    name,
    toYYYYMM(dt) AS month,
    sum(purchase)
FROM CCEvents
GROUP BY
    name,
    month
```

# #6 Awesome Table Capabilities


## Materialized Views

*Automatically aggregates data on inserts*

```
SELECT *  
FROM CCMonthlySums
```

name	month	sum(purchase)
Joe	201910	213.7000026702881
Joe	201911	296.70000553131104
Kate	201910	98.9000015258789
Kate	201911	205.5999984741211
Kate	201912	41.5

5 rows in set. Elapsed: 0.010 sec.



A benchmark performed by **Altinity** shows how fast materialized views are\*

A quote from conclusion of the benchmark: **“Using ClickHouse AggregatingMergeTree technique we decreased response time of the last point query from 4.5s to 20ms -- this is more than a 200x improvement”**

\* ClickHouse Continues to Crush Time Series, <https://www.altinity.com/blog/clickhouse-continues-to-crush-time-series>

# #6 Awesome Table Capabilities

## Live Views

*When the conditions are met, you get a instant notification*

```
SELECT *  
FROM CCEvents
```

dt	custId	name	ccNumber	posId	purchase
2019-10-11 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	557	98.9
2019-10-17 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	345	145.3
2019-10-18 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	124	17.4
2019-10-19 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	548	51
2019-11-01 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	61.1
2019-11-01 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	61.1
2019-11-13 12:00:01	123	Joe	XXXX-XXXX-XXXX-XXXX	852	45.3
2019-11-14 21:06:49	123	Joe	XXXX-XXXX-XXXX-XXXX	499	245.6
2019-11-15 21:06:49	123	Joe	XXXX-XXXX-XXXX-XXXX	756	5.8
2019-11-21 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	83.4
2019-12-07 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	20.2
2019-12-27 21:06:49	987	Kate	XXXX-XXXX-XXXX-XXXX	345	21.3

# #6 Awesome Table Capabilities

## Live Views

We want to track *Kate's spendings*, so we create a Live View

```
SELECT
    name,
    toYYYYMM(dt) AS month,
    sum(purchase) AS spending
FROM CCEvents
WHERE (name = 'Kate') AND (month = 201912)
GROUP BY
    name,
    month
```

name	month	spending
Kate	201912	41.5

```
CREATE LIVE VIEW kate_over_spending AS
SELECT
    name,
    toYYYYMM(dt) AS month,
    sum(purchase) AS spending
FROM CCEvents
WHERE (name = 'Kate') AND (month = 201912)
GROUP BY
    name,
    month
```

# #6 Awesome Table Capabilities

## Live View Tables

*When the conditions are met, you get a instant notification*

```
WATCH kate_over_spending
```

name	month	spending	_version
Kate	201912	41.5	1

```
Progress: 1.00 rows, 33.00 B (0.01 rows/s., 0.17 B/s.)
```

```
INSERT INTO CCEvents VALUES('2019-12-28 13:08:01', 987, 'Kate', 'XXXX-XXXX-XXXX-XXXX', 345, 30);
```

```
WATCH kate_over_spending
```

name	month	spending	_version
Kate	201912	41.5	1

name	month	spending	_version
Kate	201912	71.5	2

```
Progress: 2.00 rows, 66.00 B (0.01 rows/s., 0.20 B/s.)
```

# #7 Better SQL+

✓ *Formats*

✓ *Great functions:*

- *countIf, sumIf, arrayJoin, groupArray, geoLocation, JSON*

✓ *Lambda functions*

✓ *Resolving expression names*

✓ *system.numbers*

✓ *Order by condition*

Format	INSERT	SELECT
TabSeparated	✓	✓
CSV	✓	✓
Values	✓	✓
Vertical	X	✓
JSON	X	✓
TSKV	✓	✓
Protobuf	✓	✓
RowBinary	✓	✓
Null	X	✓
XML	X	✓
CapnProto	✓	X

# #7 Better SQL+

✓ *Conditional aggregates*

```
SELECT
  name,
  sum(purchase) AS spending,
  count(*) AS event_count
FROM CCEvents
GROUP BY name
```

name	spending	event_count
Kate	376	7
Joe	510.4000082015991	6

VS

```
SELECT
  sumIf(purchase, name = 'Kate') AS kate_spending,
  countIf(name, name = 'Joe') AS joe_event_count
FROM CCEvents
```

kate_spending	joe_event_count
376	6



# #7 Better SQL+

## ✓ Array functions

*arrayJoin* converts arrays to rows while *groupArray* does the opposite

```
SELECT arrayJoin([1, 2, 3]) AS arr
```

arr
1
2
3

```
SELECT groupArray(arr)  
FROM (SELECT arrayJoin([1, 2, 3]) AS arr)
```

groupArray(arr)
[1,2,3]

```
SELECT  
  arrayJoin([1, 2, 3]) AS nums,  
  arrayJoin(['a', 'b']) AS chars
```

nums	chars
1	a
1	b
2	a
2	b
3	a
3	b

Some other array functions:

- `indexOf`
- `arrayDistinct`
- `arrayReverse`
- `arrayConcat`
- `has`
- `hasAll`

... and more

# #7 Better SQL+



## Lambda functions

```
SELECT arrayMap(x -> (x * x), [1, 2, 3]) AS squared
```

squared
[1,4,9]

```
SELECT arrayFilter(x -> (x LIKE '%World%'), ['Hello', 'abc World']) AS res
```

res
['abc World']

```
SELECT arraySplit((x, y) -> y, [1, 2, 3, 4, 5], [1, 0, 0, 1, 0]) AS res
```

res
[[1,2,3],[4,5]]

*Some other functions accepting lambdas:*

- arraySum
- arrayCount
- arrayAll
- arraySort
- arrayFill

*... and more*

# #7 Better SQL+

✓ Real example for lambda functions:  
calculating **euclidean distance**

$$d(p, q) = \sqrt{\sum_{i=1}^n (q_i - p_i)^2}$$

$$d(p, q) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + (p_3 - q_3)^2}$$

```
SELECT *  
FROM embeddings
```

name	e
Joe	[0.93,0.89,0.93,...,0.11,0.58,0.47,0.81,0.95,0.4,0.4,0.76,0.8]
Joe	[0.16,0.22,0.32,...,0.17,0.9,0.91,0.3,0.69,0.77,0.63,0.52,0.5]

```
SELECT  
  name,  
  sqrt(arraySum(arrayMap((a, b) -> pow(a - b, 2), e, t))) AS distance  
FROM (  
  SELECT  
    name,  
    e,  
    (SELECT e AS target FROM embeddings WHERE name = 'Kate' ) AS t  
  FROM embeddings  
)
```

name	distance
Joe	2.5301778700047444
Kate	0

# #7 Better SQL+

✓ *Resolving expression names*

```
SELECT
    name,
    toYYYYMM(dt) AS month,
    sum(purchase) AS spending
FROM CCEvents
WHERE (name = 'Kate') AND (toYYYYMM(dt) = 201912)
GROUP BY
    name,
    toYYYYMM(dt)
HAVING sum(purchase) > 100
```

VS

```
SELECT
    name,
    toYYYYMM(dt) AS month,
    sum(purchase) AS spending
FROM CCEvents
WHERE (name = 'Kate') AND (month = 201912)
GROUP BY
    name,
    month
HAVING spending > 100
```

# #8 REST Capabilities

- ✓ *REST server with* **HTTP PORT 8123**
- ✓ *REST client with* **SELECT FROM URL**

```
$ echo 'SELECT 1 FORMAT JSON' | curl 'http://localhost:8123/' --data-binary @-
```


```
{
  "meta": [
    { "name": "1", "type": "UInt8" }
  ],
  "data": [
    { "1": 1 }
  ],
  "rows": 1,
  "statistics": {
    "elapsed": 0.000067311,
    "rows_read": 1,
    "bytes_read": 1
  }
}
```

# #8 REST Capabilities



**Secret sauce:** HTTP interface + FORMAT JSON

```
SELECT
  name,
  toYYYYMM(dt) AS month,
  sum(purchase) AS spending
FROM CCEvents
WHERE (name = 'Kate') AND
      (month = 201912)
GROUP BY
  name,
  month
FORMAT JSON
```



```
{
  "meta": [
    {"name": "name", "type": "String"},
    {"name": "month", "type": "UInt32"},
    {"name": "spending", "type": "Float64"}
  ],
  "data": [
    {"name": "Kate", "month": 201912, "spending": 71.5}
  ],
  "rows": 1,
  "statistics": {
    "elapsed": 0.002663858,
    "rows_read": 13,
    "bytes_read": 147
  }
}
```

# #8 REST Capabilities

✓ *REST client*


```
from http.server
import BaseHTTPRequestHandler, HTTPServer

class CSVHTTPServer(BaseHTTPRequestHandler):
    def do_GET(self):
        self.send_response(200)
        self.send_header('Content-type', 'text/csv')
        self.end_headers()

        self.wfile.write(bytes('Hello,1\nWorld,2\n', "utf-8"))

if __name__ == "__main__":
    server_address = ('127.0.0.1', 12345)

    HTTPServer(server_address, CSVHTTPServer).serve_forever()
```



```
CREATE TABLE rest_server (word String, value UInt64)
ENGINE=URL('http://myserver.com:12345/', CSV);
```

```
SELECT * FROM rest_server;
```

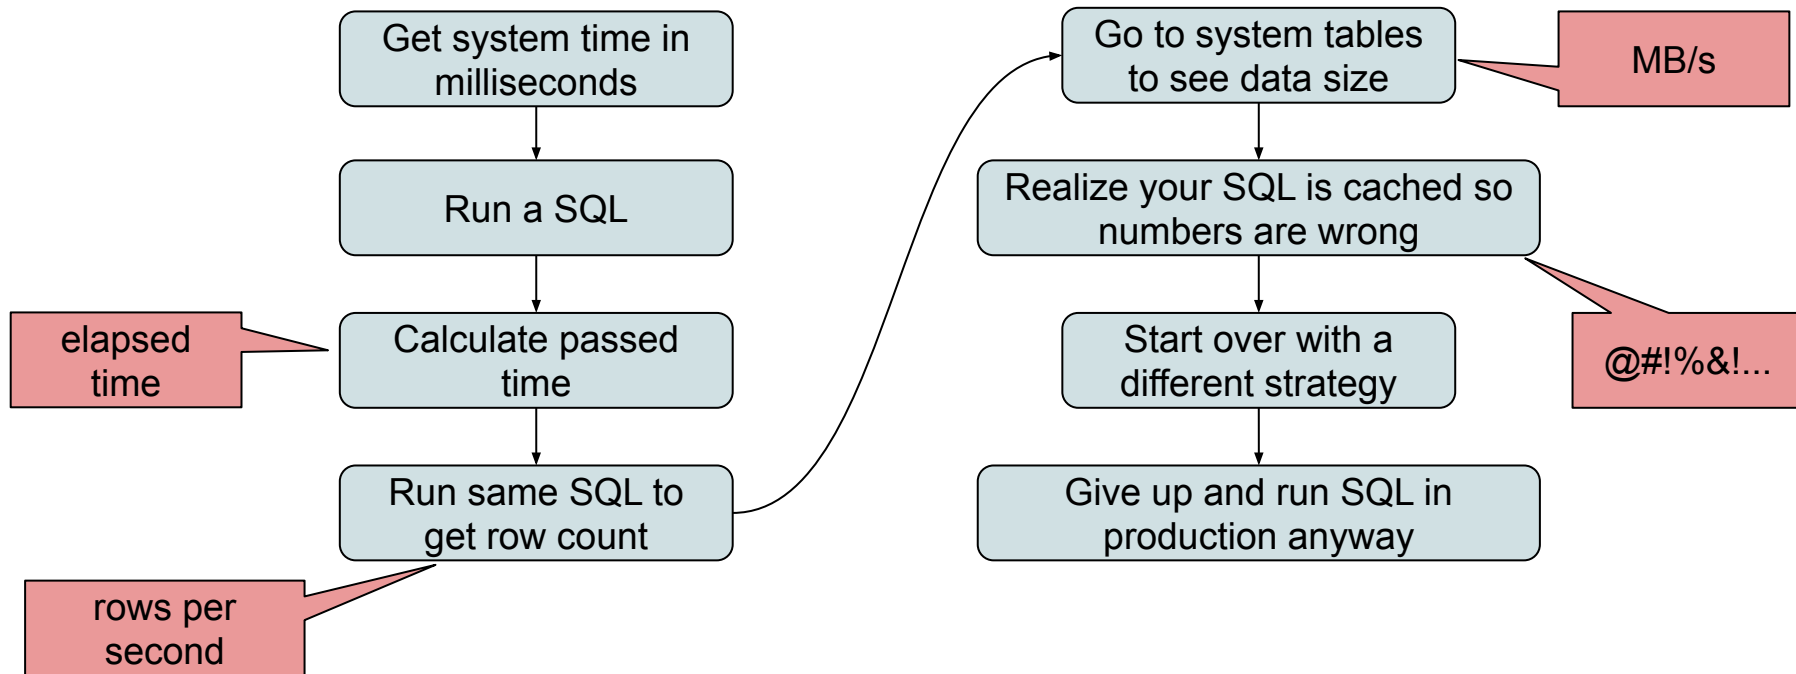
word	value
Hello	1
World	2

Use it as local table

Sample REST Server running at <http://myserver.com>

# #9 CLI

*Almost every programmer did something like this...*





# #9 CLI

✓ CLI provides a lot of useful information

```
SELECT
    table_name,
    requester,
    countIf(hit, hit = 1) AS hits,
    countIf(miss, miss = 1) AS misses,
    count() AS total,
    round((hits * 100) / total, 2) AS efficiency
FROM cachelog
GROUP BY
    table_name,
    requester
ORDER BY total DESC
LIMIT 10
```

table_name	requester	hits	misses	total	efficiency
isverenintratoken		18456146	38468	18494614	99.79
hd_v2		177647	7493254	7670901	2.32
hd_v2		20689	658119	678808	3.05
hd_v2test		77284	22	77306	99.97
isverentokentest		75086	137	75223	99.82
hd_v2		33013	41517	74530	44.29
hd_v2		9154	40209	49363	18.54
hd_v2		2998	45498	48496	6.18
hd_4aturksa		5026	8732	13758	36.53
hd_v2		9896	648	10544	93.85

elapsed  
time

rows per  
second

MB/s


10 rows in set. Elapsed: 0.711 sec. Processed 28.63 million rows, 174.83 MB (40.28 million rows/s., 246.02 MB/s.)

clickhouse :)

# #9 CLI

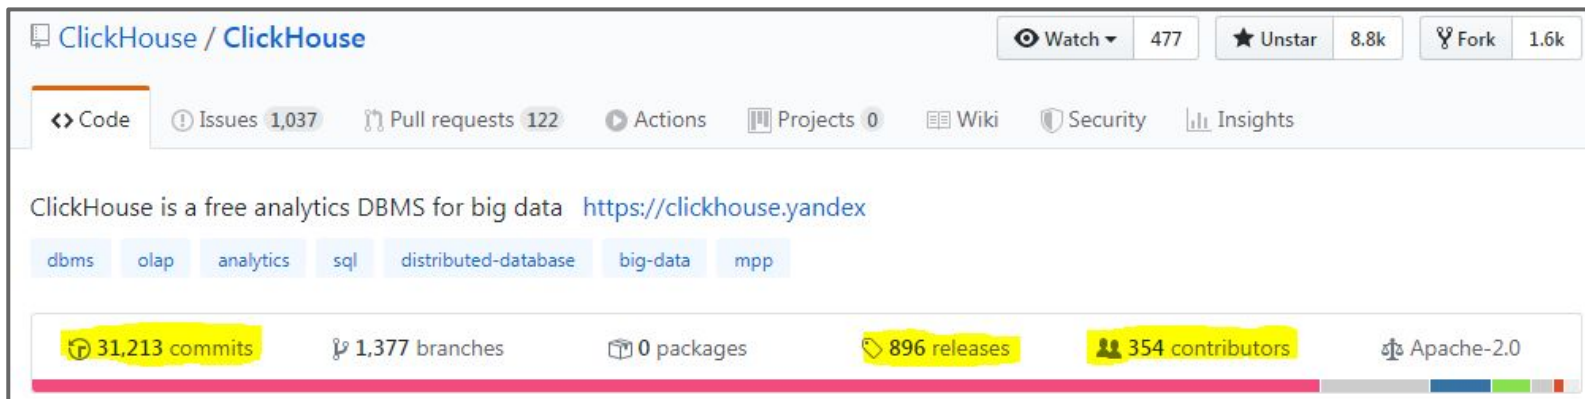
✓ *CLI provides a lot of useful information*

```
SELECT
  rand() % 100 AS percent,
  bar(percent, 1, 200, 100) AS bar
FROM system.numbers
LIMIT 10
```




# #10 Great Ecosystem

- ✓ *Open Source*
- ✓ *Free of charge!*
- ✓ *Great community support*
- ✓ *Constantly Improving*



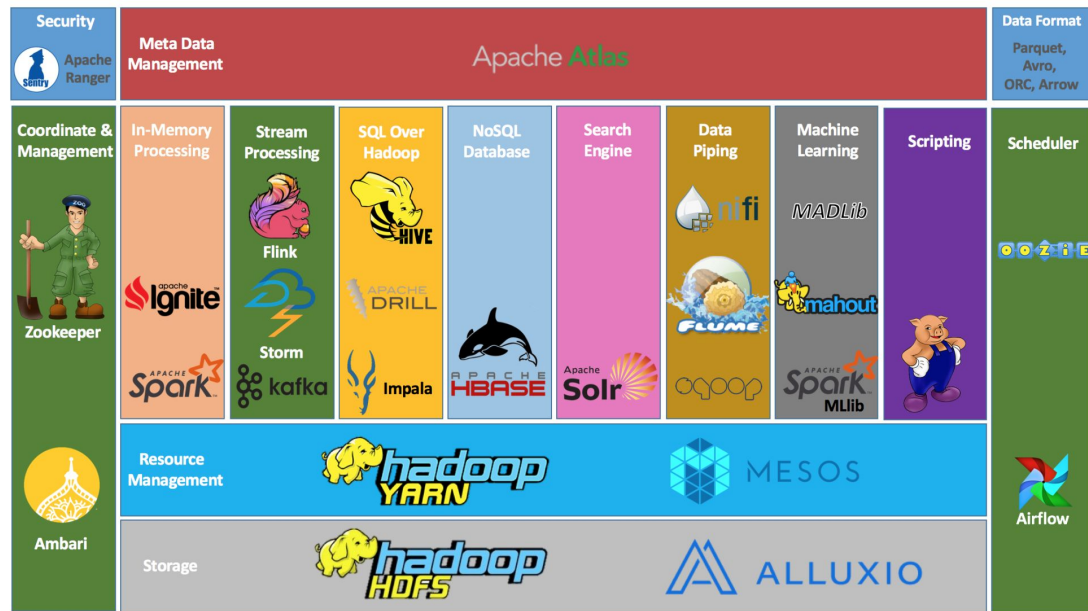
# #11 Close the gap effectively

Data is growing faster than computing power\*, so **we need to close the gap effectively**



\* <https://www.newsweek.com/2014/08/15/computers-need-be-more-human-brains-262504.html>

# #11 Close the gap effectively



Big Data Ecosystem

VS



*“Because it is easier to **learn/use one tool** that does **most of things well enough** than using **lots of tools** which are designed to do **one thing good**”*

# Honorable Mentions

- ✓ *Machine learning*
- ✓ *Data sampling*
- ✓ *Rich text functions(search, replace, RegEx, ngram)*
- ✓ *UUID, domain name and other helper functions*
- ✓ *External dictionaries*
- ✓ *Geo location*
- ✓ *clickhouse-local*
- ✓ *Distributed DDLs*

# How come ClickHouse born in Yandex and become so good?



**How come ClickHouse born in Yandex and become so good?**





**How come ClickHouse born in Yandex and become so good?**



# How come ClickHouse born in Yandex and become so good?

*ClickHouse is born because of a need at  
Yandex Metrika\**

*\* That's my opinion*

# Thank You!

Contacts:

[ramazanpolat@gmail.com](mailto:ramazanpolat@gmail.com)

<https://www.linkedin.com/in/ramazanpolat/>

You can get presentations from:

<https://github.com/ClickHouse/clickhouse-presentations>